

Assignment 2-Q3

February 4, 2026

1 MECH 309: Assignment 2, Question 3

Cagri Arslan

January 30, 2025

All work can be found on <https://github.com/imported-canuck/MECH-309>

```
[1]: # Imports
import numpy as np
from scipy import linalg
from solver import *
```

Note: I wrote all scripts by hand. Therefore, code might not be as concise/clean as a script produced by ChatGPT or taken from StackOverflow. Regardless, scripts should still be fully functional if all **assumptions** stated in docstrings are respected. I trust that this won't be penalized during grading.

a: Forward Substitution

Below is the function, `forward_sub(A, b)` that solves a linear system via forward substitution. It requires two inputs: an ($n \times n$) matrix A , assumed to be lower triangular (but not checked), and an ($n \times 1$) column vector b . It ensures that the matrix A is nonsingular and square, and that vector b is of compatible shape with A . If all of these checks pass, it builds the solution vector x : an ($n \times 1$) column vector.

```
[ ]: def forward_sub(A, b):
    """
    Solve the equation  $Ax = b$  for  $x$ , where  $A$  is a lower triangular matrix.
    Assumes a lower triangular matrix  $A$ . Assumes  $A$  is non-singular,
    and that the dimensions of  $A$  and  $b$  are compatible (but checks these and
    throws an exception if not).

    Parameters:
    A (ndarray): A lower triangular matrix of shape (n, n).
    b (ndarray): A vector of shape (n, 1).

    Returns:
    x (ndarray): The solution vector of shape (n, 1).
    """
```

```

tol = 1e-12                      # Tolerance to avoid floating point issues
n = A.shape[0]                      # Size of matrix A

if A.shape[1] != n:                # Check if A is square
    raise ValueError("Matrix A must be square.")
if b.shape[0] != n:                # Check if dimensions of b are compatible
    raise ValueError("Vector b must have compatible dimensions with matrix"
                     "A.")

x = np.zeros((n, 1))              # Initialize solution vector x as a column vector

for i in range(n):                # Loop over each row
    if np.abs(A[i, i]) <= tol:    # Check for singularity (zero diagonal
        # element)
        raise ValueError("Matrix is singular.")

    # For each column of row i before the diagonal subtract the product of
    # the entry A(i, j) and the corresponding entry of vector x from b[i]
    for j in range(i):
        b[i, 0] -= A[i, j] * x[j, 0]

    # Finally divide b[i] by the diagonal element A(i, i) to get x[i]
    x[i, 0] = b[i, 0] / A[i, i]

return x

```

Here we validate `forward_sub(A,b)` on the following matrix:

$$\begin{bmatrix} 2 & 0 & 0 \\ 3 & 1 & 0 \\ -2 & -1 & 3 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} = \begin{bmatrix} 2 \\ 6 \\ 1 \end{bmatrix} \quad (1)$$

```

[3]: A = np.array([[2.0, 0.0, 0.0],
                  [3.0, 1.0, 0.0],
                  [-2.0, -1.0, 3.0]])

b = np.array([[2.0],
              [6.0],
              [1.0]])

forward_sub(A, b)

```

```
[3]: array([[1.],
           [3.],
           [2.]])
```

b: Backward Substitution

Below is the function, `backward_sub(A, b)` that solves a linear system via backward substitution. It requires two inputs: an $(n \times n)$ matrix A , assumed to be upper triangular (but not checked), and an $(n \times 1)$ column vector b . It ensures that the matrix A is nonsingular and square, and that vector b is of compatible shape with A . If all of these checks pass, it builds the solution vector x : an $(n \times 1)$ column vector.

```
[ ]: def backward_sub(A, b):
    """
    Solve the equation  $Ax = b$  for  $x$ , where  $A$  is an upper triangular matrix.
    Assumes an upper triangular matrix  $A$ . Assumes  $A$  is non-singular,
    and that the dimensions of  $A$  and  $b$  are compatible (but checks these and
    throws an exception if not).

    Parameters:
    A (ndarray): An upper triangular matrix of shape (n, n).
    b (ndarray): A vector of shape (n, 1).

    Returns:
    x (ndarray): The solution vector of shape (n, 1).
    """
    tol = 1e-12                      # Tolerance to avoid floating point issues
    n = A.shape[0]                     # Size of matrix A
    if A.shape[1] != n:                # Check if A is square
        raise ValueError("Matrix A must be square.")
    if b.shape[0] != n:                # Check if dimensions of b are compatible
        raise ValueError("Vector b must have compatible dimensions with matrix A.")
    x = np.zeros((n, 1))               # Initialize solution vector x

    for i in range(n - 1, -1, -1):    # Loop over each row from bottom to top
        # Check for singularity (zero diagonal element)
        if np.abs(A[i, i]) <= tol:
            raise ValueError("Matrix is singular.")

        # For each column of row i after the diagonal subtract the product of
        # the entry  $A(i, j)$  and the corresponding entry of vector  $x$  from  $b[i]$ 
        for j in range(i + 1, n):
            b[i, 0] -= A[i, j] * x[j, 0]
        # Finally divide  $b[i]$  by the diagonal element  $A(i, i)$  to get  $x[i]$ 
        x[i, 0] = b[i, 0] / A[i, i]

    return x
```

Here we validate `backward_sub(A, b)` on the following matrix:

$$\begin{bmatrix} 1 & -3 & 1 \\ 0 & 2 & -2 \\ 0 & 0 & 3 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} = \begin{bmatrix} 2 \\ -2 \\ 6 \end{bmatrix} \quad (2)$$

```
[5]: A = np.array([[1.0, -3.0, 1.0],
                  [0.0, 2.0, -2.0],
                  [0.0, 0.0, 3.0]])

b = np.array([[2.0],
              [-2.0],
              [6.0]]))

backward_sub(A, b)
```

```
[5]: array([3.,
           1.,
           2.])
```

c: Gaussian Elimination

Below is the function, `gaussian_elimination(A, b)` that solves a linear system via gaussian elimination with partial pivoting. It requires two inputs: an $(n \times n)$ matrix A , and an $(n \times 1)$ column vector b . It ensures that the matrix A is nonsingular and square, and that vector b is of compatible shape with A . If all of these checks pass, it builds the solution vector x : an $(n \times 1)$ column vector. It continually applies partial pivoting to address non-singular zero-pivots and floating point issues that can occur from dividing by very small numbers.

```
[6]: def gaussian_elimination(A, b):
    """
    Solve the equation Ax = b for x using Gaussian elimination with partial
    pivoting. Expects a that A is non-singular, and that the dimensions of A
    and b are compatible (but checks these and throws an exception if not).

    Parameters:
    A (ndarray): An matrix of shape (n, n).
    b (ndarray): A vector of shape (n, 1).

    Returns:
    x (ndarray): The solution vector of shape (n, 1).
    """
    tol = 1e-12                      # Tolerance to avoid floating point issues
    n = A.shape[0]                      # Size of matrix A

    if A.shape[1] != n:                 # Check if A is square
        raise ValueError("Matrix A must be square.")
    if b.shape[0] != n:                 # Check if dimensions of b are compatible
```

```

        raise ValueError("Vector b must have compatible dimensions with matrix  $\rightarrow$ A.")

    for i in range(n - 1):           # Loop over all rows of A (apart from last,
                                    # since A is already upper-triangular by then)

        # Start partial pivoting: find the row with the largest element
        # at the column position of the pivot (don't want small numbers on pivot)
        p = i + np.argmax(np.abs(A[i:, i]))

        # If no row with a nonzero entry on the pivot exists, matrix is singular
        if np.abs(A[p, i]) <= tol:
            raise ValueError("Matrix is singular.")

        # If the current row is not the one with the greatest pivot element,
        # do partial pivoting by swapping current "row i" with row with greatest
        # pivot element "row p". And apply the same operation on vector b
        if p != i:
            A[[i, p], :] = A[[p, i], :]  # Row swap "i" and "p" on A
            b[[i, p], :] = b[[p, i], :]  # Swap corresponding entries in b

        # End partial pivoting, now eliminate entries below the pivot:
        # For each subsequent row (starting at i + 1) compute the factor that
        # would eliminate the element of row j that is below the pivot
        for j in range(i + 1, n):
            factor = A[j, i] / A[i, i]  # Compute elimination factor
            A[j] = A[j] - factor * A[i]  #  $R_j \leftarrow R_j - factor * R_i$ 
            b[j] = b[j] - factor * b[i]  # Apply the same operation on vector b

        # The matrix is now upper triangular, so x can be solved for in  $O(n^2)$ 
        # time with backward substitution (from previous part b)
        return backward_sub(A, b)

```

Here we validate gaussian_elimination(A,b) on the $n = 2$, $n = 3$, and $n = 5$ square matrices provided in solver.py

```
[7]: for CASE in ["2x2", "3x3", "5x5"]:
    print(f"\n==== Test case: {CASE} ====")

    A, b = load_test_case(CASE)

    # Make deep copies of A and b for reference solution
    A_ref = np.copy(A)
    b_ref = np.copy(b)

    x = gaussian_elimination(A, b)
    print_solution_report(A_ref, b_ref, x, label="My solver")
```

```

# Reference solution (allowed for checking)
x_ref = linalg.solve(A_ref, b_ref)
print_solution_report(A_ref, b_ref, x_ref, label="Reference (scipy.linalg.
    ↪solve)")

==== Test case: 2x2 ====
--- My solver ---
x =
[[1.]
 [2.]]
||r||_2 = 0.0

--- Reference (scipy.linalg.solve) ---
x =
[[1.]
 [2.]]
||r||_2 = 1.7763568394002505e-15

==== Test case: 3x3 ====
--- My solver ---
x =
[[1.]
 [1.]
 [1.]]
||r||_2 = 0.0

--- Reference (scipy.linalg.solve) ---
x =
[[1.]
 [1.]
 [1.]]
||r||_2 = 0.0

==== Test case: 5x5 ====
--- My solver ---
x =
[[ 1.]
 [-2.]
 [ 3.]
 [-4.]
 [ 1.]]
||r||_2 = 5.0242958677880805e-15

--- Reference (scipy.linalg.solve) ---

```

```

x =
[[ 1.]
[-2.]
[ 3.]
[-4.]
[ 1.]]
||r||_2 = 8.331852114593072e-15

```

d: LU Factorization

Below is the function, `LU_factorization(A, b)` that solves a linear system via LU factorization with partial pivoting. It requires two inputs: an $(n \times n)$ matrix A , and an $(n \times 1)$ column vector b . It ensures that the matrix A is nonsingular and square, and that vector b is of compatible shape with A . If all of these checks pass, it builds the $(n \times n)$ lower triangular matrix L and $(n \times n)$ upper triangular matrix U . It also initializes permutation matrix P as the identity matrix I and applies all row swaps to it whenever partial pivoting is triggered. It then calls performs forward substitution to compute y in $Ly = Pb$, and backward substitution to compute x in $Ux = y$. It returns matrices L , U , P and the $(n \times 1)$ solution vector x . It continually applies partial pivoting to address non-singular zero-pivots and floating point issues that can occur from dividing by very small numbers.

```
[8]: def LU_factorization(A, b):
    """
    Solve the equation Ax = b for x using LU factorization with partial
    pivoting. Expects a that A is non-singular, and that the dimensions of A
    and b are compatible (but checks these and throws an exception if not).

    Parameters:
    A (ndarray): An matrix of shape (n, n).
    b (ndarray): A vector of shape (n, 1).

    Returns:
    x (ndarray): The solution vector of shape (n, 1).
    L (ndarray): The lower triangular matrix of shape (n, n).
    U (ndarray): The upper triangular matrix of shape (n, n).
    P (ndarray): The permutation matrix of shape (n, n).
    """
    tol = 1e-12                      # Tolerance to avoid floating point issues
    n = A.shape[0]                      # Size of matrix A

    if A.shape[1] != n:                 # Check if A is square
        raise ValueError("Matrix A must be square.")
    if b.shape[0] != n:                 # Check if dimensions of b are compatible
        raise ValueError("Vector b must have compatible dimensions with matrix "
                         "A.")

    L = np.eye(n)                      # Initialize L as identity matrix
```

```

P = np.eye(n) # Initialize P as identity matrix

# Essentially apply gaussian elimination with partial pivoting to A, with
# the added step of building the lower triangular matrix L along the way.
# Whenever rows are swapped in A, the same row swap is applied to P.
for i in range(n - 1):

    p = i + np.argmax(np.abs(A[i:, i]))

    if np.abs(A[p, i]) <= tol:
        raise ValueError("Matrix is singular.")
    # If the current row is not the one with the greatest pivot element,
    # do partial pivoting by swapping row i with that row.
    if p != i:
        A[[i, p], :] = A[[p, i], :]
        # Apply same row swap to L up to column i to maintain consistency
        L[[i, p], :i] = L[[p, i], :i]
        # Apply same row swap to P
        P[[i, p], :] = P[[p, i], :]

    for j in range(i + 1, n):
        factor = A[j, i] / A[i, i]

        A[j] = A[j] - factor * A[i]
        # for Rj <- Rj - factor * Ri, insert L[j, i] = factor
        L[j, i] = factor # Store the factor in L

U = A.copy() # Upper triangular matrix is the modified A after elimination

# Solve Ly = Pb and Ux = y using O(n^2) forward and backward substitution
y = forward_sub(L, np.dot(P, b))
x = backward_sub(U, y)

return x, L, U, P

```

Here we validate LU_factorization(A,b) on the $n = 2$, $n = 3$, and $n = 5$ square matrices provided in solver.py

```

[9]: for CASE in ["2x2", "3x3", "5x5"]:
    print(f"\n==== Test case: {CASE} ====")

    A, b = load_test_case(CASE)

    # Make deep copies of A and b for reference solution
    # Note the solver modifies A and b in-place
    A_ref = np.copy(A)
    b_ref = np.copy(b)

```

```

x, L, U, P = LU_factorization(A, b)
print_solution_report(A_ref, b_ref, x, label="My solver")

# Reference solution (allowed for checking)
x_ref = linalg.solve(A_ref, b_ref)
print_solution_report(A_ref, b_ref, x_ref, label="Reference (scipy.linalg.
    ↵solve)")

```

```

==== Test case: 2x2 ====
--- My solver ---
x =
[[1.]
[2.]]
||r||_2 = 0.0

--- Reference (scipy.linalg.solve) ---
x =
[[1.]
[2.]]
||r||_2 = 1.7763568394002505e-15

==== Test case: 3x3 ====
--- My solver ---
x =
[[1.]
[1.]
[1.]]
||r||_2 = 0.0

--- Reference (scipy.linalg.solve) ---
x =
[[1.]
[1.]
[1.]]
||r||_2 = 0.0

==== Test case: 5x5 ====
--- My solver ---
x =
[[ 1.]
[-2.]
[ 3.]
[-4.]
[ 1.]]

```

```
||r||_2 = 5.0242958677880805e-15
--- Reference (scipy.linalg.solve) ---
x =
[[ 1.]
 [-2.]
 [ 3.]
 [-4.]
 [ 1.]]
||r||_2 = 8.331852114593072e-15
```